

A Case for AI to Make All of Our Decisions

One of areas I'm most interested in is Artificial Intelligence (AI). I believe that AI should make all our decisions—and I think AI *can* make all our decisions.

Decision Making Process

There's a process to making decision that we completely take for granted.

Should I order pizza or make a salad? Should I go to the gym or start a new season of *Peaky Blinders*? Do I hire the junior person willing to learn or the experienced person stuck in their ways? Should I move to Austin for the job I really want or stay here for the job I have? And should we invest more in taking care of the sick or in feeding the hungry?

Every time we make a decision to do something, we're also making a decision not to do something else. We know there are alternatives and tradeoffs associated with each of those alternatives.

For any good decision, those alternatives and tradeoffs need to be considered.

And if you think about it, we're just the consequence of a long sequence of decisions—made by us or other people. Those decisions explain how we got here, who we are, and what our potential is. Something so important can't be left so vulnerable to our impulses.

Reasons why AI Should Make All of Our Decisions

1. We need to be more objective about how we make decisions.

For an objective decision, the processes that led to the decision needs to be (1) explicit and (2) consistent with how we're making other decisions.

We can't make objective decisions. Our emotions and biases always get in the way. AI, however, ensures that these rules are met.

The obvious first question is: what about the subjective/emotional humans that are developing the AI algorithms? And yes, that's true. But we'd still be better off.

Because an AI is just a bunch of "algorithms" (or "models") layered on top of one another and those algorithm work on explicitly defined rules and principles, we'll finally need to start having some conversations about the fundamental principles at play—which is easier to separate our emotions from.

As an example: Should we enforce rent controls to help those that can't afford housing? It's obviously a very emotional and divisive topic. However, the fundamental principle at stake here is *price controls* and almost every econ textbook I've read and every economist—or even any person that's just studied some economics—I've ever talked to is clear about the effects of *price controls*: in the short term they increase demand and in the long term they decrease supply—which ultimately increase prices. And in that context, it's easy to see how if price controls were implemented for housing, even if it helped some people in the short term, in the long term it would just make housing even more out of reach for those who already can't afford it (cue the housing problem in California).

It's much easier to separate our emotions from fundamental principle discussions.

A Case for AI to Make All of Our Decisions

2. We need ways of evaluating/measuring the effectiveness and correctness of our decisions and plans to act accordingly.

There are two elements in AI called “modeling” (or “learning”) and “reinforcement learning” that force us to be explicit about how outcomes should be measured and actions that should be taken according to those results.

Have you ever heard of a politician propose a new government agency and say, “If [insert new agency] doesn’t reach x, y, or z goal in the next 5 or 10 years we’ll know we were wrong about our assumption and will backtrack accordingly”?

Our current system rewards people for not being explicit—of course no one wants to admit they were wrong.

But what if AI made the wrong decision? Of course, an AI can make wrong decisions. But these rules ensure that when wrong decisions are made, we (and of course, the AI) learns from them and course corrects when it needs to.

AI Can Make All Our Decisions

We can develop frameworks to make decisions in every field.

My field, Computer Science, already has these frameworks. We use asymptotic space and time complexity formulas to measure our coding decisions and Service Level Indicators (SLIs), like availability, throughput and response time, to measure our system design decisions.

I love studying Economics. People think Economics is the study of scarcity but it’s really the study of decision making. In a world with scarce (limited) resources (either physical resources or the time we have), what we do with those resources is what matters. There are economic models, like the Production Possibilities Frontier Model, that provide an economic value answer to how we should allocate our time. There are even models like the Lorenz Curve to help us understand what to do about income inequality. Economics, to me, is a discipline entirely focused on finding ways to measure the efficiency of our decisions on resource/time allocation.

Ethics is why I dove into Philosophy and Value Theory. Did you know philosophers like Plato, Aristotle, and Aquinas actually provided means to measure how ethical ideas are? If we want to decide whether to invest more in taking care of the sick or feeding the hungry purely on an ethical basis, we actually can.

Art used to make so uncomfortable because a framework for deciding what’s *good art* and what’s *bad art* isn’t as obvious. Until I realized that if we decide the purpose of art is to invoke an emotional response, we can measure the emotional response in the ways that it manifests. I think there are objective ways to decide how *good* a piece of art is.